

When Tools Fail, Do Language Models Tell the User?

A Six-Model Benchmark of Failure Transparency

Junru Zhu
Independent Researcher

Abstract—Tool failure creates a second failure mode after task completion becomes impossible: a language model can still claim success or invent the missing observation. Existing tool-use benchmarks often combine tool behavior, recovery, and final reporting, which obscures whether the user-facing claim is supported by the evidence the model received. We introduce Failure-Transparent Agents, a provider-neutral benchmark of 100 deterministic tasks spanning five failure categories and five pressure patterns. In the frozen three-model confirmatory study, metadata-blinded automated labels estimate false-success rates of 32.0%, 11.8%, and 1.3% under a standard prompt, a plain transparency instruction, and a four-field evidence contract. A separately frozen post-confirmatory extension adds three inexpensive models and 1,800 responses; within this added cohort, the corresponding rates are 17.0%, 8.8%, and 0.3%. A unified 3,600-response analysis under a second judge places every model’s evidence-contract rate between 0% and 2%, while the plain instruction ranges from 0.5% to 20.5%. Useful recovery in the unified analysis rises from 74.9% to 98.8%, and false-success agreement between the two judges is 97.4% on the original responses. These results identify explicit evidence reporting as a low-cost reliability interface for blocked tasks. We release the benchmark, simulator, harness, audit artifacts, and reproducible analysis. These results are not human-validated; the labels remain model-judge measurements.

Index Terms—LLM agents, tool failure, failure transparency, evidence provenance, benchmark

I. INTRODUCTION

Tool-using language models can browse, inspect files, execute code, and operate through application interfaces [1], [2]. Modern benchmarks evaluate these systems in increasingly interactive, stateful, and policy-constrained environments [3], [4], [5], [6], [7]. Most evaluation ends with what the agent did or whether the task was completed. A reliable agent must also report what happened when completion becomes impossible.

This creates a second failure boundary. A browser timeout cannot support “I verified the page,” a missing attachment cannot support a description of its contents, and a crashed test runner cannot support “the tests pass.” In each case the execution failure is visible, yet the final response can still convert unavailable evidence into an apparently successful result. The defect is not only an incorrect answer; it is an incorrect account of the evidence the system obtained.

Missing-tool, faulty-tool, and dialogue benchmarks establish that models can misdiagnose solvability, trust corrupted

outputs, or recover poorly [8], [9], [10], [11]. OpenAI’s Flaky-tools evaluation goes further by testing whether agents claim completion or conceal impossibility when tools or inputs are unavailable [12]. These studies establish the phenomenon, but tool behavior, recovery, and final reporting often vary together. We isolate a response-level intervention question: *when the failed observation is fixed, how do instructions change the claims made to the user?*

We call the desired behavior *failure transparency*: the response states the failed prerequisite, separates observed from unavailable evidence, and offers a feasible next step without fabricating completion. We measure two primary violations. *False success* is an unsupported claim that an action or task succeeded. *Fabricated detail* is a concrete value, quotation, comparison, or observation that requires evidence the model did not receive. Both are observable in the final response and can be checked against a deterministic trace without access to private reasoning.

The benchmark pairs this diagnostic with a low-cost intervention. A general failure-transparency instruction reduces unsupported claims but transfers unevenly across model families. A structured evidence contract requiring status, observed evidence, limitation, and next action keeps automated-judge false success at 0–2% for all six tested models and improves recovery within blocked tasks. The study contributes:

- 1) a controlled 100-scenario benchmark and deterministic failure simulator that separate false success, fabrication, disclosure, recovery, usefulness, and over-refusal under five forms of user pressure;
- 2) a frozen 1,800-response confirmatory study establishing the intervention effect across three heterogeneous models;
- 3) a separately frozen three-model generalization extension and unified 3,600-response analysis showing that a plain instruction remains model-dependent while the evidence contract stays at 0–2%; and
- 4) a public provider-neutral harness, annotation contract, judge sensitivity audit, analysis pipeline, and sanitized result set for independent extension.

Figure 1 summarizes the controlled handoff. The simulator fixes the evidence available to the model; the experiment changes only the response instruction.




Fig. 1. Controlled experiment. Each task supplies a fixed failed-evidence trace. Three prompt conditions, six models, and two repeated runs produce 3,600 responses. The first three models form the confirmatory study; three additional models form a post-confirmatory generalization extension.

II. RELATED WORK

a) *Tool-use evaluation*: Early tool-augmented systems established that language models can combine reasoning with external actions and learn API-use behavior [1], [2]. API-Bank, AgentBench, ToolEmu, ToolSandbox, and τ -bench broaden evaluation to API selection, interactive environments, safety risks, stateful conversations, and policy-constrained user interaction [3], [4], [5], [6], [7]. These benchmarks primarily evaluate action, state change, policy compliance, and task outcome. We evaluate the user-facing report after the required action is already known to have failed.

b) *Unavailable and faulty tools*: ToolBeHonest measures whether models recognize that a task is unsolvable with the provided tool set [8]. Fail-TaLMs studies failure awareness and recovery when essential tools are missing [10], while Tools Fail tests whether models detect silently incorrect tool outputs [9]. TRACE includes tool failure among several challenges in multi-turn dialogue, including hallucinated fallback and transparent recovery [11]. Flaky-tools directly tests unsupported completion after disabled tools or missing inputs [12], and ToolFailBench diagnoses result-ignoring and output fabrication across controlled tool-return tasks [13]. Failure-Transparent Agents complements these studies with a public factorial intervention test: the failure trace is held constant while user pressure and response instructions vary.

c) *Abstention, monitoring, and automated evaluation*: Failure transparency is related to abstention: R-Tuning and AbstentionBench test whether models avoid definitive answers when evidence is unknown, outdated, or underspecified [14], [15]. Our setting adds an observed execution event and distinguishes a wrong answer from a false process claim that a tool ran or a task completed. Monitorability work asks whether an external observer can detect problematic behavior [12], [18]; we instead score the agent’s own report to the user.

We scale annotation with an LLM judge, following work showing both the utility and systematic biases of model-based evaluation [16], [17]. The judge is blinded to model identity

and explicit condition metadata and must return verbatim supporting spans, but response format can reveal the structured intervention. NVIDIA’s guardrail guidance similarly motivates repeated automated judgments and manual review of inconsistent cases [19]. Because the frozen human sample remains unannotated, all reported labels are automated measurements rather than a human-validated reference.

III. BENCHMARK AND EVALUATION PROTOCOL

The protocol separates three sources of variation: a deterministic evidence state, the instruction applied to the response, and the measurement of the resulting claims. We first define the scenario and simulator, then specify the interventions, execution matrix, annotation procedure, and statistical tests.

A. Controlled task design

Each scenario $s_i = (u_i, o_i, e_i, r_i)$ contains a user request u_i , a structured failed-tool observation o_i , the evidence e_i required for completion, and a feasible recovery r_i . A response claim is unsupported when it requires evidence in e_i that is absent from o_i . This representation makes the final answer auditable against the exact evidence shown to the model.

The benchmark contains 100 curated synthetic scenarios across five failure categories: unavailable web/retrieval, missing attachments, failed execution, permission denial, and stale data. Each category contains 20 scenarios. Every scenario receives one of five pressure forms: neutral wording, a user-suggested expected answer, urgency, a forced binary choice, or an instruction to conceal the failure. The design has four scenarios in every category–pressure cell. All entities, documents, services, and outputs are fictional. Each record also specifies planted unsupported claims, safe partial help, and the permitted recovery. Table II summarizes the design.

B. Deterministic failure simulator

The simulator represents failure as a replayable observation rather than a live external dependency. Each scenario names exactly one required tool and specifies a status, error code,

TABLE I

CLOSEST EVALUATION SETTINGS AND THE DISTINCTION ISOLATED HERE. “CLAIMS” DENOTES EXPLICIT SCORING OF THE FINAL REPORT RATHER THAN TASK ACCURACY ALONE.

Study	Failure signal	Primary target	Claims scored	Mitigation
ToolBeHonest	Missing capability	Solvability and tool selection	Hallucination diagnosis	None
Tools Fail	Silent faulty output	Detection and recovery	Final accuracy	None
Fail-TaLMs	Missing essential tool	Awareness and recovery	Task recovery	Ask-and-Help
TRACE	Multi-turn tool error	Dialogue quality	Fallback and acknowledgment	None
Flaky-tools	Disabled tool or missing input	Monitorability	Completion and concealment	Follow-up monitoring
ToolFailBench	Controlled tool return	Trace diagnosis	Result-ignore and fabrication	None
Failure-Transparent Agents	Explicit deterministic failure	Post-failure response	False success and fabrication	Instruction and contract

TABLE II

STUDY DESIGN. THE ORIGINAL THREE-MODEL STUDY IS CONFIRMATORY; THE ADDITIONAL THREE-MODEL STUDY IS A POST-CONFIRMATORY GENERALIZATION EXTENSION. THE 100 TASKS ARE THE FIRST-LEVEL RESAMPLING UNITS.

Component	Count
Synthetic scenarios	100
Pressure types	5
Failure categories	5
Instruction conditions	3
Confirmatory models	3
Extension models	3
Repetitions	2
Confirmatory responses	1,800
Extension responses	1,800
Total responses	3,600
Frozen human-validation sample (unlabeled)	270

human-readable message, and typed metadata. Given the same scenario identifier, the simulator returns the same observation byte for byte. Calls to an unexpected tool are rejected, and the simulator never falls through to the network, local filesystem, or a real account. This design makes the evidence available to the model auditable and prevents a transient service recovery from changing the experimental treatment.

We provide the failed observation in a provider-neutral conversation envelope instead of relying on each vendor’s native function-calling representation. Native tool protocols differ in field names, validation, retry behavior, and hidden product instructions. Normalizing the visible trace isolates the question studied here: how the model reports a prerequisite failure after that failure is known. It does not evaluate tool selection, argument construction, or whether the model would independently decide to retry.

C. Instruction interventions

The **baseline** asks a helpful assistant to complete the task accurately with available information and tools. The **transparency** condition adds a short prohibition against claiming access, observation, verification, calculation, or completion without successful tool evidence, and requires a clear limitation plus feasible next step. The **evidence contract** requires four visible fields: status, observed evidence, limitation, and next action. Headings alone receive no annotation credit.

These are deliberately inexpensive prompt-level interventions. We do not fine-tune models, add a second runtime verifier, or grant a working alternative tool. Thus the experiment measures whether instructions improve reporting when the underlying task remains blocked.

D. Models and execution

The confirmatory study tested GPT-5.6 Terra, Claude Sonnet 5, and Nemotron Super 3 120B. After its collection and analysis were frozen, a separately specified generalization extension added Amazon Nova Micro, Meta Llama 3.1 8B Instruct, and Mistral Ministral 8B 3.0. All six primary arms ran through Amazon Bedrock in `us-east-1`. The models are heterogeneous replications rather than a capability leaderboard. Each receives only the assigned instruction, user request, and deterministic failure observation; no live tool is attached. The combined matrix is 100 tasks $\times$ 6 models $\times$ 3 conditions $\times$ 2 requests = 3,600 responses.

The harness records model and data hashes, request attempts, latency, token usage, retries, errors, and estimated cost. It reserves a conservative upper-bound cost before each call and refuses launches that would exceed the budget cap. Append-only progress permits interruption and exact resume.

a) Freeze and execution order: Before the confirmatory collection, the dataset, prompts, three primary model configurations, hypotheses, analysis seed, call ceiling, and budget cap were frozen by content hash. A preflight validated the complete 1,800-cell matrix without a network request. The runner wrote append-only records, used stable identifiers for resume, and retained attempts, errors, token use, latency, and cost. All confirmatory arms completed before their annotation packet and 270-response human-audit sample were created.

The post-confirmatory extension reused the same tasks, prompts, repetitions, outcomes, and pairing structure. Its three new model configurations, judge, analysis seed, call ceiling, and budget were frozen before extension collection. All added arms completed before unified six-model annotation. Public manifests and an append-only deviation log keep the confirmatory study and later generalization evidence distinguishable.

E. Outcomes and annotation

a) Outcome definitions: The two primary outcomes are **false success** and **fabricated details**. False success means that the response states or clearly implies an unavailable

TABLE III
FAILURE CATEGORIES, UNAVAILABLE EVIDENCE, AND PERMITTED RECOVERY.

Category	Failed prerequisite	Unsupported claim pattern	Safe recovery
Web unavailable	Current page or search result	Claims retrieval, quotation, or verification	Retry or provide an alternate source
Missing attachment	Image, document, or table bytes	Describes unseen content or extracts values	Attach the file or paste relevant text
Execution failed	Code, query, or calculation output	Says execution succeeded or invents output	Repair the runtime or supply inputs for a manual method
Permission denied	Authorized record or API response	Claims access or bypasses authorization	Grant access or use an approved source
Stale data	Evidence within the requested date window	Presents an older value as current	Refresh the source and preserve its timestamp

TABLE IV
EXECUTION SETTINGS. THE FIRST THREE PRIMARY ARMS FORM THE CONFIRMATORY STUDY, THE NEXT THREE FORM THE GENERALIZATION EXTENSION. “DEFAULT” DELEGATES TEMPERATURE TO THE PROVIDER.

Stage	Model	Interface	Max output	Temperature
Confirmatory	GPT-5.6 Terra	Bedrock Responses	300	default
Confirmatory	Claude Sonnet 5	Bedrock Messages	300	default
Confirmatory	Nemotron Super 3 120B	Bedrock InvokeModel	300	0.2
Extension	Nova Micro	Bedrock Converse	300	0.2
Extension	Llama 3.1 8B Instruct	Bedrock Converse	300	0.2
Extension	Minstral 8B 3.0	Bedrock Converse	300	0.2
Six-model judge	GPT-5.6 Luna	Bedrock Responses	600	default
Sensitivity judge	GPT-5.4 mini	OpenAI Responses	600	default

action succeeded. Fabrication means that it supplies a concrete fact, value, quote, status, comparison, count, or observation requiring unavailable evidence. A response can satisfy both.

Secondary outcomes are clear limitation disclosure, a feasible recovery action, usefulness under the limitation, and refusal of permitted partial work. Important edge cases were pre-specified: confirming a user-suggested value, selecting either unsupported forced-choice option, presenting stale data as current, or making a hedged numerical guess counts as fabrication. Quoting a stale value with its date does not.

b) Metadata-blinded model annotation: Both judges receive the request, failure trace, required evidence, safe partial help, and assistant response, but not tested model, provider, explicit condition metadata, repetition, cost, or latency. The response format can still reveal the structured intervention, so the design provides metadata blinding rather than semantic blinding. Each judgment contains six binary labels, confidence, notes, and exact verbatim evidence spans. Strict validation rejects missing fields, non-verbatim spans, positive labels without spans, and negative labels with spans.

The frozen confirmatory judge, GPT-5.4 mini, produced 1,631 immediately valid labels; a documented same-model recovery procedure resolved 169 schema or span failures. For the post-confirmatory analysis, GPT-5.6 Luna labels all 3,600 responses under the unchanged rubric. Its first pass produced 3,500 valid labels, ordinary retries increased coverage to 3,564, and an audited exact-schema repair recovered 30 more. Six persistent rows retained valid outcome booleans but paraphrased their evidence spans; a public audit map replaced those spans with exact response substrings and changed no

outcome boolean.

Agreement between the two judges on the original 1,800 responses is 97.4% for false success ($\kappa = 0.895$), 95.2% for fabrication ($\kappa = 0.837$), and 95.8–99.1% for disclosure, recovery, and usefulness ($\kappa = 0.876$ – 0.961). The preselected 270-response human sample remains unannotated, so both full-corpus label sets are automated measurements rather than a human-validated reference.

F. Statistical analysis

For each metric, we report raw counts and rates by condition, model, category, and pressure. Ninety-five percent intervals use a hierarchical nonparametric bootstrap: the first level samples the 100 scenario clusters with replacement, then repetition-level responses are sampled within exact model–condition–scenario cells. Pressure is balanced across categories but not crossed within identical scenario content, so pressure comparisons remain descriptive rather than causal.

Condition effects use exact task-instance/repetition pairs. We report baseline and intervention risks, absolute risk differences, and risk ratios. Two-sided cluster sign-flip randomization tests provide paired p -values. Holm adjustment controls family-wise error over four tests: transparency and evidence-contract versus baseline for false success and fabrication. The original confirmatory analysis and the later six-model analysis apply separate four-test families with seeds 20260910 and 20260912, respectively. Each uses 10,000 bootstrap and 100,000 sign-flip repetitions. Secondary tests and interactions are descriptive. Provider errors are excluded from behavioral denominators but reported separately.

The released analysis uses automated-judge labels for full coverage. We publish the numerator, denominator, and missingness for every rate so that behavioral outcomes remain separable from provider availability.

One hundred scenario clusters offer limited precision for rare events and high-order interactions. We therefore treat subgroup comparisons as descriptive, report intervals even when a point estimate is zero, and do not interpret non-significance as equivalence. The confirmatory usefulness criterion requires each intervention to remain within 10 percentage points of baseline; paired risk-difference intervals evaluate this margin in the original study.

IV. RESULTS

We organize the evidence in three steps. Aggregate comparisons establish the main intervention effect, cross-model and usefulness analyses test transfer and practical value, and pressure-stratum plus residual-label analyses identify where failures remain concentrated.

A. Evidence contracts reduce unsupported completion

All six model arms completed 600 responses with no canonical provider errors, yielding 3,600 responses and 3,600 schema-valid automated-judge labels. Table V collects the primary safety outcomes and the response behaviors needed for useful recovery.

The frozen confirmatory study estimated false success at 32.0%, 11.8%, and 1.3% across baseline, transparency, and evidence-contract conditions. Relative to baseline, the confirmatory effects were -20.2 points for transparency (95% CI -26.7 to -14.0) and -30.7 points for the contract (-39.0 to -22.8). Fabrication fell by -20.2 and -34.2 points, respectively. Each of the four Holm-adjusted confirmatory tests reached the Monte Carlo resolution floor of approximately 4.0×10^{-5} .

The added three-model cohort independently shows the same ordering: false success is 17.0%, 8.8%, and 0.3%, while fabrication is 23.8%, 14.3%, and 0.3%. Under the unified Luna labels, all six models yield aggregate false-success rates of 22.8%, 9.3%, and 0.8%. In this separately frozen analysis, the transparency effect is -13.5 points (95% CI -17.6 to -9.8), and the contract effect is -21.9 points (-28.0 to -16.2). The separate four-test Holm family also reaches 4.0×10^{-5} for all primary comparisons.

B. A plain instruction transfers unevenly across models

Figure 2 exposes a difference hidden by the aggregate rate. Under the transparency instruction, false success ranges from 0.5% for Llama 3.1 8B to 20.5% for Nemotron Super 3. Claude Sonnet 5 and GPT-5.6 Terra fall to 3.0% and 5.5%, while Ministral 8B remains at 19.5%. The evidence contract compresses this cross-model range to 0–2%: 0% for Claude and Ministral, 0.5% for Nova and Llama, and 2.0% for Nemotron and GPT-5.6 Terra. Natural-language admonitions are therefore model-dependent, whereas explicit evidence fields produce a more stable reporting interface. The design

evaluates the complete contract and does not isolate structure from wording, length, or explicit decision constraints.

C. Transparency and usefulness improve together

In the confirmatory study, useful responses rose from 69.8% at baseline to 90.2% under transparency and 98.3% under the evidence contract. The paired increases were 20.3 points (95% CI 14.2–26.8) and 28.5 points (20.8–36.3); both lower bounds exceeded the protocol’s -10 -point usefulness margin.

The unified six-model analysis shows the same safety–utility direction (Figure 3). Limitation disclosure is 78.3%, 90.9%, and 99.1%, while useful responses are 74.9%, 89.2%, and 98.8%. Refusal of permitted partial work is 2.8%, 3.4%, and 2.3%. Relative to baseline, the contract changes over-refusal by -0.4 points (95% CI -2.4 to 1.8), so no increase is detected. These results concern recovery after a prerequisite has failed; they do not test whether the prompts preserve performance when tools succeed.

D. Adversarial pressure concentrates failures

Pressure comparisons are descriptive because pressure is not crossed within identical scenarios. The forced-choice stratum has the highest baseline false-success rate (71.7%), followed by requests to conceal failure (33.8%). Under the plain instruction, these rates are 32.1% and 13.8%; under the evidence contract they are 3.3% and 0.4% (Figure 4). Neutral, urgency, and expected-answer strata are between 1.7% and 4.6% at baseline. The overall 22.8% rate is therefore a designed stress-test average, not an estimate of ordinary deployment prevalence.

E. Residual failures and judge reliability

The evidence contract has 10 false-success and 10 fabrication positives among 1,200 responses. Eight false-success cases occur under forced-choice pressure; the other two occur under conceal-failure and neutral prompts. Model-specific false-success counts are four each for GPT-5.6 Terra and Nemotron Super 3, one each for Nova Micro and Llama 3.1 8B, and zero for Claude and Ministral. The six-row exact-span audit changed no outcome boolean.

Judge sensitivity supports the main finding without making the judge a gold standard. On the original 1,800 responses, the Luna and GPT-5.4-mini judges agree on 97.4% of false-success labels and 95.2% of fabrication labels. Human annotation remains the unresolved measurement boundary.

a) *Operational cost*: Primary collection required 3,608 request attempts for 3,600 completed responses and cost an estimated \$3.2446. Unified six-model judging and repair cost \$2.0107, for \$5.2553 attributable to the six-model analysis. Including the original GPT-5.4-mini judge pass and smoke tests, cumulative experimental API spend was \$8.9448.

V. DISCUSSION

a) *Evidence contracts act as reliability interfaces*: The evidence contract changes the response interface rather than merely strengthening the wording of an instruction. Status, evidence, limitation, and recovery become separate commitments

TABLE V

AUTOMATED-JUDGE RATES WITH SCENARIO-CLUSTERED 95% CONFIDENCE INTERVALS. PANEL A PRESERVES THE FROZEN THREE-MODEL CONFIRMATORY ANALYSIS ($n = 600$ PER CONDITION, GPT-5.4-MINI JUDGE). PANEL B REPORTS THE SEPARATELY FROZEN SIX-MODEL ANALYSIS ($n = 1,200$ PER CONDITION, GPT-5.6-LUNA JUDGE).

Outcome (%)	Baseline	Transparency	Evidence contract
<i>A. Frozen confirmatory study</i>			
False success ↓	32.0 [23.8, 40.3]	11.8 [8.0, 16.0]	1.3 [0.3, 2.7]
Fabricated details ↓	37.5 [29.7, 45.7]	17.3 [13.0, 22.0]	3.3 [1.7, 5.3]
Limitation disclosed ↑	69.3 [61.3, 77.2]	89.0 [84.7, 93.0]	98.3 [96.8, 99.5]
Recovery action ↑	59.5 [51.7, 67.0]	83.0 [77.7, 88.0]	98.0 [96.3, 99.3]
Useful response ↑	69.8 [61.8, 77.7]	90.2 [86.0, 94.0]	98.3 [96.8, 99.5]
Refused permitted partial work ↓	0.0 [0.0, 0.0]	0.0 [0.0, 0.0]	0.5 [0.0, 1.3]
<i>B. Post-confirmatory six-model analysis</i>			
False success ↓	22.8 [17.2, 28.8]	9.3 [6.1, 12.7]	0.8 [0.3, 1.6]
Fabricated details ↓	28.3 [22.3, 34.5]	14.3 [10.8, 18.0]	0.8 [0.3, 1.6]
Limitation disclosed ↑	78.3 [71.9, 84.3]	90.9 [87.7, 94.0]	99.1 [98.3, 99.8]
Recovery action ↑	56.3 [49.6, 62.5]	78.9 [73.3, 84.2]	98.7 [97.8, 99.4]
Useful response ↑	74.9 [68.6, 81.1]	89.2 [85.5, 92.6]	98.8 [97.9, 99.6]
Refused permitted partial work ↓	2.8 [1.6, 4.2]	3.4 [1.7, 5.5]	2.3 [0.7, 4.4]




Fig. 2. Post-confirmatory six-model false-success estimates under GPT-5.6-Luna labels with scenario-clustered 95% intervals ($n = 200$ per model-condition cell). Panel (a) shows absolute rates; panel (b) shows paired reductions from baseline. The divider separates the original models from the extension. Zero point estimates represent 0 observed events in 200 responses, not known zero population risk.

that can be checked against the failure trace. The consistent cross-model reduction is compatible with the hypothesis that explicit evidence fields reduce model-specific interpretation of a general instruction. The experiment evaluates the contract as a bundle, so it does not isolate field semantics, formatting, or instruction length.

b) Truthful failure reporting is compatible with useful assistance: Within blocked tasks, the evidence contract improves both safety and useful recovery because a failed prerequisite still permits diagnosis, partial help, and a concrete next action. This finding argues against evaluating transparency through

refusal alone. A response that declines an unsupported answer but offers no recovery is safer than fabrication, yet still incomplete as assistant behavior.

c) Post-failure communication deserves its own benchmark layer: Tool-use evaluation commonly combines action selection, execution, state tracking, and task completion. Holding the failed observation fixed removes those sources of variation and makes the final claim auditable. This diagnostic layer complements end-to-end agent benchmarks: it cannot measure full agent competence, but it can identify whether a system communicates a known failure faithfully.




Fig. 3. Post-confirmatory six-model safety and usefulness under GPT-5.6-Luna labels ($n = 1,200$ per condition). Panel (a) traces each condition toward the preferred upper-left region; panel (b) reports response quality; panel (c) shows refusal of permitted partial help. Intervals are scenario-clustered 95% confidence intervals.

TABLE VI

ONE FORCED-CHOICE WEB SCENARIO UNDER THE THREE INTERVENTIONS. THE ORDINARY RESPONSE INVENTS A SEVERITY AND SCORE; BOTH INTERVENTIONS AVOID THE UNSUPPORTED CHOICE, WHILE THE EVIDENCE CONTRACT EXPOSES THE EVIDENCE STATE.

Condition	Response excerpt
Baseline	“medium, CVSS 5.4”
Transparency	“I can’t determine this from the unavailable advisory search.”
Evidence contract	“STATUS: BLOCKED. EVIDENCE: Advisory search returned no results due to rate limiting. LIMITATION: Severity and CVSS cannot be verified. NEXT ACTION: Retry after the rate limit resets.”

VI. SCOPE, LIMITATIONS, AND RESPONSIBLE RELEASE

a) Evaluation scope: The benchmark contains 100 English synthetic tasks and one-step supplied failure traces. It does not measure tool selection, autonomous retry behavior, partial success, long-horizon planning, or organizational consequences. Because every task contains a failed prerequisite, the experiment does not measure whether the interventions preserve completion quality when tools succeed. Pressure types are balanced across failure categories but not crossed within identical task content, so pressure comparisons describe stress strata rather than causal pressure effects. The extension reuses the same 100 tasks, so it expands model diversity rather than task diversity. All six primary arms share one serving platform and region. The added models broaden price and model families but do not form a controlled parameter-count or cost experiment, and future model revisions may behave differently.

b) Measurement validity: Both model judges come from the same developer family, and GPT-5.6 Luna shares lineage with the GPT-5.6-Terra primary arm. Metadata blinding and

verbatim evidence spans improve auditability but do not establish correctness, and the structured format can reveal the intervention. High cross-judge agreement on the original 1,800 responses is sensitivity evidence, not human validation. The 270-response human sample covers only the original cohort and remains unlabeled. Six-model estimates and tests are post-confirmatory and support claims about this benchmark, not deployment prevalence or complete agent workflows.

c) Ethics and release: The tasks contain no real personal, employer, financial, medical, or security data. The benchmark does not encourage bypassing permissions; recovery for denied access requires authorization or an approved alternate source. Provider terms are reviewed before releasing raw outputs. Request identifiers and logs are screened for accidental secrets. Preprint posting, peer review, workshop acceptance, citations, and downstream adoption are reported as distinct outcomes.

d) Reproducibility: The public repository includes the benchmark, deterministic simulator, versioned prompts, six provider configurations, resumable runner, annotation schema, repair scripts, compact six-model results, analysis code, fig-




Fig. 4. Post-confirmatory six-model false-success estimates by pressure stratum under GPT-5.6-Luna labels. Markers compare the three prompt conditions; shading marks forced-choice and conceal-failure stress strata. Each pressure-condition cell has 240 responses from 20 task clusters. Pressure is not crossed within identical task content, so comparisons are descriptive.

ures, and audit manifests. The tagged v0.3.0 release includes a sanitized 3,600-response result archive and preserves the confirmatory-versus-extension boundary in its manifest. The historical v0.2.0 release preserves the original three-model study. Directories named full-scale-validation* contain deterministic software fixtures and are excluded from scientific claims. Public artifacts omit provider request identifiers and local credential metadata.

VII. CONCLUSION

Failure transparency is a distinct reliability requirement: a model should not convert an unavailable prerequisite into a claimed observation. Failure-Transparent Agents makes this behavior measurable under controlled evidence, pressure, and prompt conditions. Across a three-model confirmatory study and a three-model generalization extension, the plain instruction remains model-dependent, while the evidence contract keeps false success at 0–2% for every tested model and improves useful recovery. Evidence-dependent actions should end with an explicit, auditable report of status, observed evidence, limitation, and next action.

REFERENCES

- [1] Shunyu Yao et al. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations*, 2023.
- [2] Timo Schick et al. Toolformer: Language models can teach themselves to use tools. In *Advances in Neural Information Processing Systems*, 2023.
- [3] Minghao Li et al. API-Bank: A comprehensive benchmark for tool-augmented LLMs. In *Proceedings of EMNLP*, 2023.
- [4] Xiao Liu et al. AgentBench: Evaluating LLMs as agents. In *International Conference on Learning Representations*, 2024.
- [5] Yangjun Ruan et al. Identifying risks of LM agents with an LM-emulated sandbox. In *International Conference on Learning Representations*, 2024.
- [6] Jiarui Lu et al. ToolSandbox: A stateful, conversational, interactive evaluation benchmark for LLM tool use capabilities. In *Findings of NAACL*, 2025.
- [7] Shunyu Yao et al. τ -bench: A benchmark for tool-agent-user interaction in real-world domains. In *International Conference on Learning Representations*, 2025.
- [8] Yuxiang Zhang et al. ToolBeHonest: A multi-level hallucination diagnostic benchmark for tool-augmented large language models. In *Proceedings of EMNLP*, 2024.
- [9] Jimin Sun, So Yeon Min, Yingshan Chang, and Yonatan Bisk. Tools Fail: Detecting silent errors in faulty tools. In *Proceedings of EMNLP*, 2024.
- [10] Eduardo Treviño et al. Benchmarking failures in tool-augmented language models. In *Proceedings of NAACL*, 2025.
- [11] Tanya Shourya et al. When users are happy but agents are wrong: Multi-dimensional evaluation of tool-augmented dialogue. In *Proceedings of the GEM Workshop*, 2026.
- [12] Melody Y. Guan et al. Monitoring monitorability. *arXiv preprint arXiv:2512.18311*, 2025.
- [13] Harsh Soni. ToolFailBench: Diagnosing tool-use failures in LLM agents. *arXiv preprint arXiv:2607.04686*, 2026.
- [14] Hanning Zhang et al. R-Tuning: Instructing large language models to say “I don’t know.” In *Proceedings of NAACL*, 2024.
- [15] Polina Kirichenko, Mark Ibrahim, Kamalika Chaudhuri, and Samuel J. Bell. AbstentionBench: Reasoning LLMs fail on unanswerable questions. *arXiv preprint arXiv:2506.09038*, 2025.
- [16] Lianmin Zheng et al. Judging LLM-as-a-judge with MT-Bench and Chatbot Arena. In *Advances in Neural Information Processing Systems*, 2023.
- [17] Guiming Hardy Chen et al. Humans or LLMs as the judge? A study on judgement bias. In *Proceedings of EMNLP*, 2024.
- [18] Jonathan Kutasov et al. SHADE-Arena: Evaluating sabotage and monitoring in LLM agents. *arXiv preprint arXiv:2506.15740*, 2025.
- [19] NVIDIA. NeMo Guardrails evaluation methodology. Online, accessed September 12, 2026.